

When Aerodynamic States Become Latent: Reserve-Risk Screening under Degraded Wind-Turbine SCADA Telemetry

Junyu Li and Juntao Du* School of Statistics and Applied Mathematics, Anhui University of Finance and Economics, Bengbu 233030, China

*Corresponding author: dujuntao@aufe.edu.cn

Abstract—Operating reserve screening near the demarcation boundary between Maximum Power Point Tracking (MPPT, Region 2) and active blade-pitch regulation (Region 3) is governed by acute asymmetric shortfall risk: generation over-forecasts incur severe balancing penalties relative to surplus energy headroom. Deterministic aerodynamic power curves presume fresh, fully observable telemetry; however, commercial SCADA systems routinely encounter transmission latency, packet serialization dropouts, and aggregator-boundary pitch unobservability. Under degraded telemetry, the turbine control register is unobservable at deployment, while its electromechanical consequences remain partially observable in SCADA feeds ($Z_t \rightarrow P_t, Q_t$, trajectory). Across multi-seed benchmarks on the 134-turbine WTB commercial plant and multi-year records from three European facilities, this paper addresses a single governing question: *what is the minimum sufficient modeling strategy under changing SCADA observability regimes?*

Our findings resolve three operational information regimes: (1) Fresh and observable $\rightarrow$ Physics: Under pristine telemetry, deterministic aerodynamic rules achieve the lowest reserve-screening surrogate cost in the aerodynamic boundary band ($|v_{\text{phys}} - 10.5| \leq 1.0$ m/s, $N = 13,883$ events) among all evaluated approaches (589,535 kW · h at 10-minute dispatch $h = 1$ and 881,367 kW · h at 1-hour dispatch $h = 6$, with $\sim 6.8\%$ violation), outperforming deep neural networks by 8.1%. (2) Stale but observable $\rightarrow$ Recalibration: While a 60-minute latency contingency serves as a synthetic stress endpoint where unadapted physical rules undergo reliability breakdown ($24.0\% \pm 1.7\%$ violation at $h = 1$, leaving 127.6 MWh of unhedged shortfall), 10–30 minute delays demonstrate that when telemetry channels remain observable, non-neural state-conditional recalibration absorbs 55.3% of this shortage loss (127.6 $\rightarrow$ 57.0 MWh) without parameter updates. (3) Critical state unobservable $\rightarrow$ Consequence representation: When blade-pitch registers are withheld at deployment, learned representations (trained using historically available pitch information that is withheld at deployment) reconstruct latent operating regimes from secondary electromechanical and spatio-temporal consequences, doubling transition-window recall (0.416 vs. 0.196). Relative to the unadapted physical comparator under pitch withholding, learned representations reduce the reserve-screening surrogate by 46.7k–68.6k kW · h across evaluated latencies; however, recoverability does not imply decision superiority, as strong wind-speed-conditioned quantiles remain lower-cost plant-wide. Dynamic Mixture-of-Experts routing confers no statistical advantage over unrouted dense baselines ($p = 0.380$), while decoupled modular architectures match tail reliability (9.7% violation). Furthermore, heuristic selective decision abstention fails to outperform random rejection, being Pareto-dominated by uniform reserve margin expansion. These findings establish that machine learning is needed when observability fails, not whenever prediction is difficult.

Index Terms—Wind turbine operating reserve; SCADA telemetry degradation; partial observability; latent operating boundary; transmission latency; aerodynamic power curve; reliability breakdown; reserve-risk screening; state-conditional recalibration; MPPT-to-pitch transition.

I. INTRODUCTION

In bulk power system operations, real-time generation shortfalls incur severe asymmetric shortage penalties relative to surplus headroom; conventional symmetric metrics like root-mean-square error (RMSE) mask severe tail shortfall events that govern operating reserve adequacy in Level-1 pre-dispatch screening [1], [2], [3], [4], [5]. This operational vulnerability concentrates near the aerodynamic demarcation boundary between Maximum Power Point Tracking (MPPT, Region 2) and active blade-pitch regulation (Region 3). Near rated inflow velocity, local power sensitivity transitions abruptly from a steep cubic trajectory ($\partial P / \partial v \propto v^2$) toward zero as pitch actuators shed aerodynamic lift [6], [7]. When operating states shift across this threshold unannounced, deterministic power curves over-extrapolate generation into Region 3, precipitating severe unhedged reserve deficits.

This vulnerability is exacerbated by telemetry degradation in commercial Supervisory Control and Data Acquisition (SCADA) systems, including gateway serialization overhead, packet queuing backlogs, and intermittent cellular transmission jitter [8], [9]. In operational archives (such as Kelmarsh), 99.6% of discrete turbine status transitions occur between standard 10-minute polling boundaries. Furthermore, in third-party aggregator and Virtual Power Plant (VPP) triage, blade-pitch registers are frequently restricted across commercial protocol firewalls. When delayed or missing measurements enter physical power equations, dispatch systems clear reserve margins against obsolete operating points. Under degraded telemetry, the turbine operating regime becomes a partially observed latent state: the control register itself is unobservable at deployment, while its electromechanical consequences remain partially observable in arrived SCADA feeds ($Z_t \rightarrow P_t, Q_t$, trajectory).

Existing literature frequently treats aerodynamic curves and deep neural networks as competing predictive paradigms [10], [11], [12], [13], [14], [15], [16], overlooking their distinct operational boundaries under information freshness constraints. Rather than asking which model achieves marginal predictive

gains, bulk power system operations require answering a foundational engineering question:

What is the minimum sufficient model
under changing telemetry observability? (1)

To resolve this question, we structure our investigation across three distinct information regimes, bounded by an operational non-equivalence principle:

- **Fresh and observable** $\rightarrow$ **Physics**: Under pristine SCADA telemetry, deterministic aerodynamic quantile curves achieve the lowest reserve-screening surrogate cost within the active boundary band ($|v_{\text{phys}} - 10.5| \leq 1.0$ m/s, $N = 13,883$ cell events) among all evaluated methods (589,535 kW · h at $h = 1$ and 881,367 kW · h at $h = 6$, with target-satisfying $\sim 6.8\%$ violation), saving 8.1% (−77,659 kW · h) over neural architectures.
- **Stale but observable** $\rightarrow$ **Recalibration**: Operational delays of 10–30 minutes exhibit distribution drift where non-neural state-conditional recalibration absorbs 55.3% of shortage loss (127.6 $\rightarrow$ 57.0 MWh) without updating network weights. A 60-minute latency contingency serves strictly as an asymptotic reliability breakdown stress endpoint ($24.0\% \pm 1.7\%$ violation at $h = 1$) rather than a claim of normal industrial prevalence.
- **Critical state unobservable** $\rightarrow$ **Consequence representation**: When blade pitch is withheld, neural representations infer latent operating regimes from electromechanical consequences ($Z_t \rightarrow P_t, Q_t$, trajectory), doubling transition recall (0.416 vs. 0.196) and reducing reserve screening penalties by 46.7k–68.6k kW · h over unadapted physical curves.
- **Boundary Condition (Recoverability $\nRightarrow$ Decision superiority)**: Reconstructing the hidden boundary does not yield universal decision superiority: strong wind-speed-conditioned quantiles remain lower-cost plant-wide (+523,044 kW · h posterior penalty), while posterior conditioning operates as a localized risk-hedging mechanism that contracts violation and shortage exposure during operating transitions at the expense of additional reserve procurement.

The 134-turbine WTB site serves as the primary mechanism-identification environment (with five random seeds quantifying training stochasticity rather than independent physical replications), while multi-year evaluations on three external European facilities (ENGIE La Haute Borne, Kelmarsh, Penmanshiel) probe transferability and delineate site-dependent boundary conditions. Treating spatio-temporal graph and mixture-of-experts formulations as established diagnostic baselines, this investigation establishes four verifiable contributions:

1. **Regime 1 — Aerodynamic Sufficiency and Reliability Breakdown**: We establish that deterministic physical rules are the minimum sufficient and lowest-cost model under fresh observable telemetry in the boundary band (589,535 kW · h at $h = 1$, 881,367 kW · h at $h = 6$). Operational latencies of 10–30 min expose the onset of cubic extrapolation errors, while a 60-min contingency lag

demonstrates asymptotic reliability breakdown ($24.0\% \pm 1.7\%$ violation, 127.6 MWh deficit).

2. **Regime 2 — Sufficiency of Non-Neural Recalibration under Observable Drift**: We decouple telemetry staleness from genuine channel missingness, demonstrating that under full channel observability, non-neural state-conditional recalibration absorbs 55.3% of shortage loss (127.6 $\rightarrow$ 57.0 MWh) without parameter updates.
3. **Regime 3 — Consequence-Based State Recovery and Decision Boundary**: When blade pitch is withheld at deployment, learned representations recover the latent operating regime from secondary electromechanical consequences ($Z_t \rightarrow P_t, Q_t$, trajectory). We prove that recoverability does not imply plant-wide decision superiority: strong wind-speed-conditioned quantiles remain lower-cost plant-wide, while posterior conditioning serves as a localized risk-hedging mechanism during regime transitions.
4. **Falsification of Alternative Explanations and Cross-Site Heterogeneity**: Dynamic Mixture-of-Experts routing confers no statistical advantage over unrouted dense baselines ($p = 0.380$), while decoupled modular architectures match tail reliability (9.7% violation). External evaluations across four commercial wind farms reveal heterogeneous information regimes rather than uniform transfer superiority, where the observed pattern is consistent with exploitable spatial redundancy as a possible moderator (not turbine count alone), while turbine count alone does not explain the variation.

II. RELATED WORK

A. *From Average Wind-Power Accuracy to Transition-Window Risk*
Spatio-temporal graph neural networks (STGNNs) advance multi-step wind plant forecasting by modeling spatial cross-correlations and wake advection [10], [11], [12], [13], with physics-informed extensions incorporating wake and boundary constraints [14], [15], [16]. While these architectures reduce plant-wide L_2 error, continuous trajectory minimization obscures discrete turbine operating state transitions between MPPT (Region 2) and pitch regulation (Region 3) that govern spinning reserve adequacy [17], [18], [19].

Recent physics-aware dynamic-graph mixture-of-experts (MoE) architectures, such as that of Zhao et al. [20], test whether coupling spatial learning, physical guidance, and expert routing improves point forecasting accuracy. Rather than iterating on forecasting network topologies, our formulation isolates the upstream information-and-decision problem: when critical aerodynamic states become unobservable due to telemetry degradation, what regime information remains recoverable from secondary consequence channels, when non-neural conditional recalibration suffices without parameter updates, and whether the recovered latent state alters asymmetric reserve-shortfall decisions. Probabilistic forecasts translate predictive uncertainty into reserve allocation [1], [2], [21]. We address this operational triage problem: screening aerodynamic transition risks under degraded SCADA telemetry prior to market clearing, functioning strictly as a Level-1 pre-dispatch proxy while abstracting downstream transmission constraints.

III. PROBLEM FORMULATION AND LATENT OPERATING-BOUNDARY FRAMEWORK

A. Level-1 Pre-Dispatch Risk Screening and Asymmetric Decision Problem (PSREI)

Operating reserve screening functions as a pre-dispatch risk assessment mechanism to mitigate costly real-time generation imbalances [21], [22].

Hierarchical grid dispatch interface: Our formulation is strictly bounded to Level 1 pre-dispatch operating reserve screening proxies at the plant EMS or aggregator terminal, abstracting Level 2 bulk transmission AC-OPF and wholesale balancing settlements.

Penalized Reserve-Shortfall Energy Index (PSREI): For scheduled forecast $\hat{y}_{i,t}$ and realized generation $y_{i,t}$, shortfall is $s_{i,t} = \max(\hat{y}_{i,t} - y_{i,t}, 0)$. For upward reserve r_b and penalty ratio ρ , the **Penalized Reserve-Shortfall Energy Index (PSREI)** in kW · h is:

$$C(r_b; \rho) = \sum_{(i,t)} [r_b(i,t) + \rho \max(\hat{y}_{i,t} - y_{i,t} - r_b(i,t), 0)] \Delta t. \quad (2)$$

Differentiating $\mathbb{E}[C(r_b; \rho)]$ yields the first-order condition $\partial \mathbb{E}[C]/\partial r_b = 1 - \rho \Pr[s_{i,t} > r_b] = 0$, giving the Newsvendor critical fractile $q^*(\rho) = 1 - 1/\rho$ [1]. For $\rho = 10$, $q^* = 0.90$, establishing a nominal 10% violation target.

Loss geometry and pinball quantile head: Point forecasting minimizes symmetric L_2 loss converging to conditional mean $\mathbb{E}[y | \mathbf{x}]$, which departs from upper quantiles in asymmetric operations ($\rho = 10$). Upward reserve margin $\hat{r}_{i,t}$ covering shortfall $s_{i,t} = \max(\hat{y}_{i,t} - y_{i,t}, 0)$ is estimated via pinball loss:

$$\mathcal{L}_q(s, \hat{r}) = (s - \hat{r})(q - \mathbf{1}[s < \hat{r}]) = \begin{cases} q(s - \hat{r}), & s \geq \hat{r}, \\ (1 - q)(\hat{r} - s), & s < \hat{r}. \end{cases} \quad (3)$$

Setting $q = q^* = 0.90$ aligns training with the Newsvendor condition: Bayes-risk-minimizing prediction converges to conditional shortfall quantile $\hat{r}^* = Q_{0.90}(s | \mathbf{x})$ [1], [2].

PCC portfolio smoothing & dispatch horizons: Net power at the Point of Common Coupling (PCC) bus aggregates turbine outputs ($P_{\text{farm},t} = \sum_{i=1}^M P_{i,t}$, $R_{\text{farm},t} = \sum_{i=1}^M R_{i,t}$), where turbulence cancels across the array, leaving boundary transitions as dominant unhedged risk. We evaluate immediate dispatch ($h = 1$, 10-min ahead) dominated by serial autocorrelation, and operational dispatch ($h = 6$, 1-hour ahead) where spatial wake dynamics govern decisions.

B. Latent Operating Regimes: MPPT, Pitch Regulation, and Demarcation Boundaries

Let $G = (V, E)$ denote a turbine array spatial graph with N nodes. For each turbine $i \in V$ at time step t , the system observes feature vector $\mathbf{x}_{i,t} \in \mathbb{R}^F$ and predicts active power $y_{i,t} \in \mathbb{R}$. Given observation history H and horizon P , the forecast mapping is:

$$\hat{\mathbf{Y}}_{t+1:t+P} = \mathcal{F}(\mathbf{X}_{t-H+1:t}, \mathcal{A}_{t-H+1:t}), \quad (4)$$

where $\mathbf{X}_{t-H+1:t} \in \mathbb{R}^{H \times N \times F}$ and $\mathcal{A}_{t-H+1:t}$ denotes the graph adjacency sequence. We enforce strict causal boundaries: all input features, graph adjacency structures, and supervisory anchors depend solely on telemetry available at or before dispatch anchor t , with target trajectories commencing strictly at

$t+1$. In WTB, instantaneous active power $\text{Patv}_{i,t}$ serves solely as an anchor-time operational status indicator, preventing look-ahead leakage.

At dispatch time t , turbine i operates in a discrete aerodynamic control regime governed by inflow wind speed $w_{i,t} = \text{Wspd}_{i,t}$ and collective blade pitch angle $\bar{p}_{i,t} = \frac{1}{3} \sum_{k=1}^3 p_{i,t}^{(k)}$:

$$Z_{i,t} = \begin{cases} 0, & w_{i,t} < u_{\text{idle}} \quad (\text{idle}), \\ 1, & u_{\text{idle}} \leq w_{i,t} \leq u_{\text{rated}}, \bar{p}_{i,t} < p_{\text{th}} \quad (\text{MPPT}), \\ 2, & w_{i,t} > u_{\text{rated}}, \bar{p}_{i,t} \geq p_{\text{th}} \quad (\text{pitch-regulated}), \\ 3, & \text{otherwise} \quad (\text{transitional/ambiguous}). \end{cases} \quad (5)$$

Thresholds follow turbine specifications: WTB (134 turbines: $u_{\text{idle}} = 3.0$ m/s, $u_{\text{rated}} = 10.5$ m/s, $p_{\text{th}} = 2.0^\circ$), Kelmars (6 MM92: 3.0 m/s, 12.5 m/s, 1.0°), Penmanshiel (14 operational MM82, numbered WT01–WT15 excluding WT03: 3.0 m/s, 14.5 m/s, 1.0°), and ENGIE La Haute Borne (4 M-M82: 3.0 m/s, 14.5 m/s, 1.0°). Supervised alignment masks ambiguous transitions ($Z_{i,t} = 3$) via $M_{i,t} = \mathbf{1}[Z_{i,t} \neq 3]$.

Under pristine telemetry, $\bar{p}_{i,t}$ and $w_{i,t}$ are known, so $Z_{i,t}$ is directly observable. Under transmission latency $\tau > 0$ or when blade-pitch registers are withheld across aggregator interfaces, the control register itself ceases to be accessible at deployment, while its dynamical electromechanical consequences ($Z_t \rightarrow P_t, Q_t$, trajectory) remain partially preserved in arrived SCADA feeds. Hence, $Z_{i,t}$ is a **control-state latent variable** rather than an unobservable physical process. The central inferential object is the soft operating-boundary posterior:

$$p(Z_{i,t} | \mathbf{X}_{\leq t}), \quad (6)$$

inferred from historical trajectory, active-power transients, and spatial wake consequence channels.

C. Aerodynamic Power Conversion and the Stale-State Control Boundary

Mechanical power captured from inflow air is governed by $P_{\text{mech}}(t) = \frac{1}{2} \rho_{\text{air}} \pi R^2 C_p(\lambda(t), \beta(t)) v^3(t)$, where $\rho_{\text{air}} \approx 1.225$ kg/m³, R is rotor radius, $v(t)$ is wind speed, $\beta(t)$ is pitch angle, and $\lambda(t) = \omega_r(t)R/v(t)$ is tip-speed ratio [6], [7]. In Region 2 (MPPT, $u_{\text{idle}} \leq v \leq u_{\text{rated}}$), blades maintain nominal fine pitch $\beta \approx 0^\circ$ yielding cubic power growth $P(v) \propto v^3$; in Region 3 (Pitch regulation, $u_{\text{rated}} < v \leq u_{\text{cut-out}}$), the pitch controller rotates blades to shed lift, clamping electrical output at nameplate capacity P_{rated} :

$$\left. \frac{\partial P}{\partial v} \right|_{v \rightarrow u_{\text{rated}}^-} = \frac{3}{2} \rho_{\text{air}} \pi R^2 C_{p,\text{max}} u_{\text{rated}}^2 \gg 0, \quad \left. \frac{\partial P}{\partial v} \right|_{v \rightarrow u_{\text{rated}}^+} \approx 0. \quad (7)$$

When SCADA telemetry experiences latency $\tau > 0$, aerodynamic curves evaluate stale tuples $(v_{t-\tau}, \beta_{t-\tau})$. If the turbine enters Region 3 during the latency window ($Z_{t-\tau} = 1$ while $Z_t = 2$), the delayed pitch register remains at fine pitch $\beta_{t-\tau} \approx 0^\circ$. Consequently, the aerodynamic rule extrapolates power along the cubic trajectory:

$$\hat{P}_{\text{phys}}(t) = \frac{1}{2} \rho_{\text{air}} \pi R^2 C_p(\lambda(t), 0^\circ) v^3(t) \gg P_{\text{rated}}. \quad (8)$$

Because output clamps at P_{rated} , this stale evaluation produces severe over-forecast error $\hat{P}_{\text{phys}}(t) - P_{\text{rated}} \gg 0$, surging empirical violation rates to $24.0\% \pm 1.7\%$ under Delay-6

($\rho = 10$), which acts as an asymptotic reliability breakdown stress endpoint.

D. The Non-Neural Baseline: State-Conditional Quantile Recalibration

To evaluate whether latency-induced drift can be absorbed by non-neural statistical correction, we formalize state-conditional quantile recalibration. Given delayed telemetry $\mathbf{s}_{i,t-\tau} = (w_{i,t-\tau}, \bar{p}_{i,t-\tau})$, we partition the state space into K bins $\{\mathcal{B}_k\}_{k=1}^K$ defined by cut-in, rated wind speed, and pitch thresholds ($Z_{i,t-\tau} = k$). To eliminate data leakage, empirical quantile reserve margins $\hat{r}(k)$ are calibrated strictly on held-out validation residuals $\mathcal{D}_k^{\text{val}}$:

$$\hat{r}(k) = \inf \left\{ r \in \mathbb{R} : \frac{1}{|\mathcal{D}_k^{\text{val}}|} \sum_{(i,t) \in \mathcal{D}_k^{\text{val}}} \mathbf{1}[\hat{y}_{i,t}^{\text{phys}} - y_{i,t} \leq r] \geq q^* \right\}, \quad (9)$$

where $\mathcal{D}_k^{\text{val}} = \{(i, t) \in \mathcal{D}_{\text{val}} : \mathbf{s}_{i,t-\tau} \in \mathcal{B}_k\}$ and $q^* = 1 - 1/\rho = 0.90$.

We evaluate two calibration protocols: (1) *Clean-Frozen Calibration*, where margins $\hat{r}^{\text{clean}}(k)$ are calibrated on clean validation data ($\tau = 0$) and strictly frozen; and (2) *Condition-Matched Calibration*, where for a known delay τ , margins $\hat{r}(k; \tau)$ are re-estimated on held-out validation residuals with matching delay τ . Recalibration is an **information-preserving correction**: it updates residual margins conditional on observed bins, but cannot synthesize missing state information when blade pitch is withheld.

E. Latent State Recovery from Spatio-Temporal Consequence Channels

Physical dynamic wake graph & encoder: For WTB, time-varying directed wake graphs $\mathcal{A}_t$ link turbine pairs within aerodynamic wake cones ($\alpha = 25^\circ$) weighted by streamwise and cross-stream decay [14], [16]. A compact directed-diffusion GRU encoder (110k parameters, 3.8 ms single-window execution) aggregates messages via:

$$\mathbf{X}_t^{(\ell+1)} = \text{LN}(\mathbf{W}_{\text{self}}^{(\ell)} \mathbf{X}_t^{(\ell)} + \mathbf{W}_{\text{in}}^{(\ell)} \text{Agg}_{\text{in}}(\mathbf{X}_t^{(\ell)}, \mathcal{A}_t) + \mathbf{W}_{\text{out}}^{(\ell)} \text{Agg}_{\text{out}}(\mathbf{X}_t^{(\ell)}, \mathcal{A}_t) + \mathbf{R}^{(\ell)} \mathbf{X}_t^{(\ell)}), \quad (10)$$

and extracts node-level representations $\mathbf{h}_{i,t}$. When blade pitch $\bar{p}_{i,t}$ is unobservable, the encoder infers latent regime $p(Z_{i,t} | \mathbf{X}_{\leq t})$ from secondary consequence channels: active power transients ($\text{Patv}_{i,t}$), electrical responses, and dynamic wake advection, strictly excluding $\text{Pab_mean}_{i,t}$.

F. Reserve-Tail Estimation and Residual Quantile Sizing

To shield operations from balancing exposure, reserve margins target the asymmetric shortfall distribution rather than symmetric point forecasts, separating: (1) posterior operating-state inference $\pi_{i,t}^{\text{pitch}} = p(Z_{i,t} = 2 | \mathbf{X}_{\leq t})$; (2) scheduled point forecasting $\hat{\mathbf{y}}_{i,t+1:t+P}$; and (3) reserve-tail quantile sizing $\hat{r}_{i,t} = Q_{0.90}(s_{i,t} | \mathbf{X}_{\leq t})$.

Directly estimating extreme quantiles on raw power trajectories via uncalibrated neural pinball regression introduces severe tail conservatism ($>1.6\text{M kW}$ reserve, $<1.2\%$ violation, Table~I). To resolve this pathology, estimation is decoupled: point trajectories $\hat{\mathbf{y}}$ anchor nominal expectations, while upward reserve margins $\hat{r}_{i,t}$ are estimated conditionally on forecast

residuals:

$$\hat{r}_{i,t} = \text{MLP}_{\text{res}}([\mathbf{h}_{i,t}; \mathbf{x}_{\text{anchor}}]), \quad (11)$$

minimized under $\mathcal{L}_{q^*}(s_{i,t}, \hat{r}_{i,t})$,

or binned empirically over held-out validation residuals conditioned on latent posterior $\pi_{i,t}^{\text{pitch}}$.

G. Architecture Variants as Mechanism Ablations

To test whether dynamic Mixture-of-Experts (MoE) routing provides operational value over unrouted representations, we formulate three matched-capacity variants: 1. **STGQ-Modular** (Decoupled Canonical): Pairs frozen spatio-temporal representations $\mathbf{h}_{i,t}$ with a dedicated residual quantile head; 2. **STGQ-Dense** (Unrouted Joint): Couples spatial embeddings $\mathbf{h}_{i,t}$ directly to a shared linear projection for point forecasts $\hat{\mathbf{y}}$ and a joint dense quantile head; 3. **STGQ-Routed** (Dynamic MoE): Evaluates dynamic MoE routing [23], [24] combining expert predictions $\hat{\mathbf{y}}_{i,t+1:t+P} = \sum_{e=1}^E g_{i,t}^{(e)} f_e(\mathbf{h}_{i,t})$.

Training minimizes:

$$\mathcal{L} = \mathcal{L}_{\text{pred}} + \lambda_{\text{align}} \mathcal{L}_{\text{align}} + \lambda_{\text{aux}} \mathcal{L}_{\text{aux}} + \lambda_{\text{smooth}} \mathcal{L}_{\text{smooth}} + \lambda_{\text{bal}} \mathcal{L}_{\text{bal}} + \lambda_{\text{force}} \mathcal{L}_{\text{force}}, \quad (12)$$

where $\mathcal{L}_{\text{align}}$ anchors logits to declared regimes $Z_{i,t}$ on unmasked samples ($M_{i,t} = 1$); $\mathcal{L}_{\text{force}}$ sharpens boundary discrimination; $\mathcal{L}_{\text{bal}}$ balances expert utilization; and $\mathcal{L}_{\text{aux}}, \mathcal{L}_{\text{smooth}}$ enforce spatial wake consistency.

Explicit Methodological Notice: $\mathcal{L}_{\text{align}}$ acts strictly as an **operating-regime prior** embedded into representation space. The neural backpropagation pipeline contains no quantile loss, pinball penalty, or Newsvendor decision loss. Checkpoint selection is strictly governed by validation RMSE. Quantile reserve margins are calibrated on held-out validation residuals or estimated by the decoupled residual head, yielding a jointly represented boundary posterior rather than an end-to-end decision posterior.

IV. CASE STUDY CONFIGURATION AND OPERATIONAL CONSTRAINTS

A. Datasets and Preprocessing

Empirical evaluation spans the 134-turbine WTB benchmark and the ERA5 thermodynamic contrast patch. WTB (KDD Cup 2022) comprises 134 turbines over 245 days of 10-minute SCADA telemetry ($T = 35,280$, $N = 134$, $F = 11$) [25], with chronological 180/30/35-day train/val/test splits ($H = 36$, $P = 24$). ERA5 covers three archived months on a 16×16 grid ($T = 2208$, $N = 256$, 8 features) [26] under 1325/441/442 splits. Multi-year SCADA records from three European commercial wind plants (ENGIE La Haute Borne, Kelmars, Penmanshiel) provide diverse turbine architectures for cross-farm evaluation.

B. Controlled Telemetry Degradation Protocols and Realism Mapping

To evaluate resilience under communication stress, we simulate three degradation tiers across five seeds (201–205): (1) **Tier 1 (Nominal Polling, 0–10 min, Clean)**: fresh SCADA on 10-min grid (IEC 61400-25); (2) **Tier 2 (Plausible Degraded Telemetry, 10–30 min, Delay-1 to 3)**: transmission lags from substation buffer backlogs [9], characterizing the operational transition regime; and (3) **Tier 3 (Contingency Stress Envelope, 60 min, Delay-6)**: a 6-step synthetic lag evaluating asymptotic physical breakdown [27], [28] as a stress endpoint rather than an industrial prevalence claim. Perturbations in-




Fig. 1. Conceptual architecture of the minimum sufficient modeling principle across three telemetry observability regimes. Incoming SCADA feeds enter an Observability & Staleness Gate that routes decisions to the least complex sufficient model: Tier 1 routes fresh observable telemetry ($\tau = 0, \exists\beta$) to deterministic aerodynamic rules ([Lowest Cost], 6.8% violation); Tier 2 routes observable operational delays ($\tau \leq 20$ min, $\exists\beta$) to state-conditional recalibration, absorbing 55.3% of shortage loss without neural retraining; Tier 3 routes pitch-withheld or contingency telemetry ($\beta = \emptyset$ or $\tau \geq 60$ min) to learned consequence representations (STGQ-Modular), recovering latent regimes while preserving the central boundary condition that recoverability does not imply plant-wide decision superiority.

clude additive Gaussian noise ($\sigma_v = 1.0$ m/s, $\sigma_\beta = 2.0^\circ$) and a two-state Markov burst dropout chain ($p_{GB} = 0.08, p_{BB} = 0.75, d \leq 6$) [9].

Fig.~1 illustrates the causal decision failure mechanism and the three-tier operational recovery hierarchy across telemetry degradation regimes.

C. Pre-Dispatch Risk Screening Framework and Resampling Hierarchy

All neural architectures share an identical training protocol using AdamW with early stopping on validation RMSE, gradient clipping, and mixed-precision execution across five random seeds (201–205). Strong baselines (Graph WaveNet, Graph Transformer, GAT-GRU, PatchTST, iTransformer, TiDE) repeat across seeds 201–205 under identical evaluation caches. To eliminate data leakage, the sequential dataflow $\mathcal{D}_{\text{train}} \rightarrow \mathcal{D}_{\text{val}} \rightarrow \text{Freeze} \rightarrow \mathcal{D}_{\text{test}}$ freezes quantile offsets $\hat{r}(k)$ prior to test replay. We report 5-seed initialization statistics alongside operational 95% confidence intervals computed via a paired daily cluster bootstrap over the 35 observed test days (144 ten-minute steps per daily cluster), using 1,000 Monte Carlo bootstrap resamples (where resampled replicates reflect bootstrap draws rather than independent observational units).

V. EMPIRICAL FINDINGS: MINIMUM SUFFICIENT MODELING ACROSS OBSERVABILITY REGIMES

A. Regime 1: Fresh Observable Telemetry $\rightarrow$ Deterministic Aerodynamics and Operational Breakdown Boundaries (RQ1)

Table~I summarizes the multi-day replay dispatch benchmark for 1-hour ahead operational dispatch ($h = 6$) evaluated on the boundary-active operating slice ($|v_{\text{phys}} - 10.5| \leq 1.0$ m/s, $N = 13,883$ cell observations across 134 turbines over the 35-

day test split) under nominal clean telemetry. Table~II reports seed-paired hypothesis tests and mechanism ablations across operational regimes, along with 95% confidence intervals from a paired daily cluster bootstrap over the 35 observed test days (1,000 resamples).

The empirical evidence directly establishes the operational sufficiency of deterministic aerodynamics under fresh observable telemetry. Under pristine telemetry ($\tau = 0$), Continuous Physical Quantile achieves the lowest surrogate reserve screening cost within the active boundary band ($|v_{\text{phys}} - 10.5| \leq 1.0$ m/s, $N = 13,883$ cell observations) among all evaluated methods ($589,535 \pm 77,390$ kW $\cdot$ h at immediate dispatch $h = 1$ and $881,367 \pm 95,741$ kW $\cdot$ h at 1-hour dispatch $h = 6$, with target-satisfying $\sim 6.8\%$ violation). It outperforms both unrouted dense networks (STGQ-Dense: 991,181 kW $\cdot$ h) and dynamically routed architectures (STGQ-Routed: 959,027 kW $\cdot$ h) by 8.1%–11.1%. As confirmed in Table~II, Continuous Physical Quantile achieves a $-77,659$ kW $\cdot$ h advantage over neural representations, with a 95% daily cluster bootstrap confidence interval (1,000 paired resamples over 35 observed days) strictly excluding zero ($[-106,175, -49,143]$). When inflow wind speed and blade pitch angle are fresh and directly observable, deterministic aerodynamic equations yield the lowest evaluated surrogate reserve screening cost. Neural networks, lacking hard physical saturation bounds, introduce unnecessary variance under nominal telemetry.

However, Table~I exposes the operational reliability boundary of physical rules when telemetry becomes stale. In real SCADA networks, transmission delays of 10–30 minutes introduce latency drift between true inflow and evaluated registers. Under a 60-minute latency contingency stress test (Delay-6), which serves as an asymptotic reliability breakdown stress endpoint rather than an industrial prevalence claim, physical equations evaluate stale state tuples ($v_{t-\tau}, \beta_{t-\tau}$). Near rated wind speed, the localized slope discontinuity (7) amplifies delayed-inflow errors: turbines that have crossed into active blade-pitch regulation ($Z_t = 2$) are evaluated with obsolete fine pitch ($\beta_{t-\tau} \approx 0^\circ$), over-extrapolating power output along the steep cubic trajectory (8). Consequently, unadapted physical quantile violation rates surge from 6.81% to $24.02\% \pm 1.69\%$ at $h = 1$ and $12.20\% \pm 1.25\%$ at $h = 6$ (1,136,466 kW $\cdot$ h), breaching the nominal 10% Newsvendor target ($q^* = 0.90$) and leaving 127.55 MWh of unhedged generation shortfall across the array.

B. Regime 2: Stale but Observable Telemetry $\rightarrow$ Non-Neural Recalibration Absorbs Drift (RQ2)

When telemetry channels remain observable under transmission delays, condition-matched state-conditional recalibration (9) adjusts quantile margins using held-out validation residuals undergoing the identical delay. As shown in Table~III, state-conditional recalibration absorbs 55.3% of the shortage loss ($127.55 \rightarrow 57.01$ MWh), restoring fleet violation to 9.46% and achieving a surrogate reserve cost of 1,228,609 kW $\cdot$ h at $h = 1$. Under full channel observability, recalibrated physical curves outperform all neural models (STGQ-Routed: 1,274,047 kW $\cdot$ h; STGQ-Modular: 1,283,855 kW $\cdot$ h).

TABLE I

CROSS-SEED MULTI-REGIME PRE-DISPATCH RESERVE SCREENING BENCHMARK ON BOUNDARY-ACTIVE OPERATING REGIMES UNDER CLEAN-CALIBRATED PROTOCOL ($h = 6$, 1-HOUR AHEAD DISPATCH, 5 SEEDS 201–205). MEAN $\pm$ STANDARD DEVIATION ACROSS BOUNDARY-BAND EVENTS ($|v_{\text{phys}} - 10.5| \leq 1.0$ m/s, $N = 13,883$ CELL OBSERVATIONS, 35 TEST DAYS, 134 TURBINES, $\Delta t = 10$ MIN, COST RATIO $\rho = 10$). ALL QUANTILE BINS AND MULTIPLIERS CALIBRATED EXCLUSIVELY ON NOMINAL CLEAN VALIDATION DATA. COMMENSURATE PLANT-WIDE FULL-POPULATION BASELINES ($N = 408,939$ CELLS) RANGE FROM 14.0M TO 16.5M kW · h (SEE TEXT AND METRIC-SCOPE REGISTRY).

Regime	Evaluated Model	Penalized Reserve Energy (kW·h)	Violation Rate	Reserve (kW)	Shortage (kW·h)	Pinball Loss
Clean (boundary)	Continuous Physical Quantile	881,367 $\pm$ 95,741	6.81% $\pm$ 1.13%	653,738 $\pm$ 72,557	22,763 $\pm$ 4,013	31.41 $\pm$ 2.66
	Direct Quantile MLP [‡]	1,649,721 $\pm$ 285,682	0.58% $\pm$ 0.26%	1,630,572 $\pm$ 293,346	1,915 $\pm$ 958	64.62 $\pm$ 12.05
	Global Quantile	998,988 $\pm$ 107,360	6.46% $\pm$ 1.31%	771,304 $\pm$ 95,661	22,768 $\pm$ 4,854	36.50 $\pm$ 3.41
	STGQ-Dense	991,181 $\pm$ 110,211	6.54% $\pm$ 1.04%	761,480 $\pm$ 94,244	22,970 $\pm$ 4,121	36.16 $\pm$ 3.58
	STGQ-Routed	959,027 $\pm$ 103,290	7.10% $\pm$ 1.17%	707,249 $\pm$ 91,305	25,178 $\pm$ 4,816	34.77 $\pm$ 3.16
	Missingness-Aware GBDT	1,238,603 $\pm$ 112,051	4.50% $\pm$ 0.73%	1,085,168 $\pm$ 118,015	15,343 $\pm$ 2870	46.85 $\pm$ 3.48
Delay-6	Continuous Physical Quantile	1,136,466 $\pm$ 84,101	12.20% $\pm$ 1.25% [†]	662,335 $\pm$ 73,483	47,413 $\pm$ 3,091	38.85 $\pm$ 2.06
	Direct Quantile MLP [‡]	1,682,251 $\pm$ 284,323	1.20% $\pm$ 0.52%	1,651,979 $\pm$ 297,715	3,027 $\pm$ 1,594	62.15 $\pm$ 11.70
	Global Quantile	1,180,603 $\pm$ 91,754	10.85% $\pm$ 1.55% [†]	780,916 $\pm$ 96,853	39,969 $\pm$ 6,087	40.73 $\pm$ 2.43
	STGQ-Dense	1,179,859 $\pm$ 95,268	11.04% $\pm$ 1.22% [†]	773,103 $\pm$ 94,462	40,676 $\pm$ 4,788	40.70 $\pm$ 2.66
	STGQ-Routed	1,163,593 $\pm$ 91,299	11.82% $\pm$ 1.64% [†]	720,015 $\pm$ 92,196	44,358 $\pm$ 7,341	40.01 $\pm$ 2.54
	Missingness-Aware GBDT	1,320,003 $\pm$ 93,135	8.36% $\pm$ 0.97%	1,030,858 $\pm$ 114,266	28,915 $\pm$ 4,469	46.68 $\pm$ 2.45
Noise	Continuous Physical Quantile	998,273 $\pm$ 106,253	7.59% $\pm$ 0.95%	746,697 $\pm$ 82,474	25,158 $\pm$ 3,968	33.97 $\pm$ 2.78
	Direct Quantile MLP [‡]	1,639,909 $\pm$ 242,922	0.80% $\pm$ 0.28%	1,618,439 $\pm$ 249,374	2,147 $\pm$ 953	60.62 $\pm$ 9.85
	Global Quantile	1,053,552 $\pm$ 114,150	7.48% $\pm$ 1.48%	802,361 $\pm$ 99,513	25,119 $\pm$ 5,822	36.26 $\pm$ 3.24
	STGQ-Dense	1,048,083 $\pm$ 115,181	7.63% $\pm$ 1.13%	792,472 $\pm$ 92,391	25,561 $\pm$ 4,643	36.04 $\pm$ 3.28
	STGQ-Routed	1,027,077 $\pm$ 114,317	7.62% $\pm$ 1.09%	769,864 $\pm$ 106,672	25,721 $\pm$ 3,963	35.16 $\pm$ 3.19
	Missingness-Aware GBDT	1,314,640 $\pm$ 129,158	4.66% $\pm$ 0.80%	1,160,329 $\pm$ 137,650	15,431 $\pm$ 3,099	47.11 $\pm$ 3.87
Markov	Continuous Physical Quantile	907,871 $\pm$ 96,820	7.20% $\pm$ 1.15%	663,875 $\pm$ 73,821	24,400 $\pm$ 4,106	31.59 $\pm$ 2.61
	Direct Quantile MLP [‡]	1,680,339 $\pm$ 289,602	0.57% $\pm$ 0.25%	1,660,135 $\pm$ 298,104	2,020 $\pm$ 992	64.38 $\pm$ 12.00
	Global Quantile	1,030,131 $\pm$ 108,592	6.87% $\pm$ 1.39%	785,305 $\pm$ 97,397	24,483 $\pm$ 5,104	36.78 $\pm$ 3.35
	STGQ-Dense	1,020,078 $\pm$ 112,195	6.90% $\pm$ 1.07%	775,899 $\pm$ 96,235	24,418 $\pm$ 4,246	36.35 $\pm$ 3.57
	STGQ-Routed	990,625 $\pm$ 104,487	7.51% $\pm$ 1.22%	720,858 $\pm$ 92,754	26,977 $\pm$ 5,183	35.10 $\pm$ 3.10
	Missingness-Aware GBDT	1,261,941 $\pm$ 112,754	4.88% $\pm$ 0.75%	1,094,730 $\pm$ 120,612	16,721 $\pm$ 3,298	46.62 $\pm$ 3.40

[†]Exceeds nominal 10% violation target ($q^* = 0.90$). Bold indicates lowest surrogate metric. **Scope & Reconciliation:** Evaluated at $h = 6$ strictly across boundary-band cells ($|v_{\text{phys}} - 10.5| \leq 1.0$ m/s, $N = 13,883$, $\sim 2.4\%$ of test split); Continuous Physical Quantile achieves lowest cost (881,367 kW · h) but collapses to 12.20% violation under Delay-6. Full-population ($N = 408,939$) direct quantiles achieve 14.0M–16.5M kW · h at $h = 6$ and 6.49M–8.05M at $h = 1$ (Section V-D; incommensurate scopes). [‡]Frozen-Embedding Direct Quantile MLP evaluates uncalibrated pinball regression on frozen representations; calibrated STGQ-Modular resolves this conservatism.

TABLE II

SEED-PAIRED OPERATIONAL HYPOTHESIS TESTS AND MECHANISM ABLATION CONTRASTS ($h = 6$, 1-HOUR AHEAD DISPATCH, COST RATIO $\rho = 10$).

Regime	Evaluated Contrast	Δ PSREI (kW·h)	95% Daily Cluster CI (1,000 Resamples)
Clean	Cont. Phys. (Rule)	−77,659	[−106.2k, −49.1k]**
	STGQ-Dense vs. Routed	+32,154	[−1.1k, +65.5k]
	Global Quantile	+39,962	[+12.6k, +67.4k]*
	Missingness GBDT	+279,576	[+249.5k, +309.6k]*
	Direct MLP (Pinball)	+690,695	[+448.8k, +932.6k]**
Delay-6	Cont. Phys. (Breakdown)	−27,126	[−46.9k, −7.3k] [†]
	STGQ-Dense vs. Routed	+16,267	[+7.1k, +25.4k]*
	Global Quantile	+17,010	[+3.5k, +30.5k]*
	Missingness GBDT	+156,411	[+140.0k, +172.8k]*
	Direct MLP (Pinball)	+518,658	[+307.2k, +730.1k]**
Noise	Cont. Phys. (Robust)	−28,804	[−47.6k, −10.1k]*
	STGQ-Dense vs. Routed	+21,006	[+2.4k, +39.6k]*
	Global Quantile	+26,475	[+9.1k, +43.8k]*
	Missingness GBDT	+287,563	[+248.8k, +326.3k]*
	Direct MLP (Pinball)	+612,832	[+413.7k, +811.9k]**
Markov	Cont. Phys. (Rule)	−82,754	[−111.0k, −54.5k]**
	STGQ-Dense vs. Routed	+29,453	[−2.9k, +61.8k]
	Global Quantile	+39,506	[+16.0k, +63.1k]*
	Missingness GBDT	+271,316	[+245.0k, +297.6k]*
	Direct MLP (Pinball)	+689,714	[+448.7k, +930.7k]**

Note: 95% CIs from paired daily cluster bootstrap over 35 observed test days (144 steps/cluster, 1,000 Monte Carlo resamples).

** $p < 0.01$, * $p < 0.05$. [†]Physical rule breaches nominal 10% violation target (12.20%). Dynamic MoE routing advantage over unrouted dense architecture is falsified ($p = 0.380$).

To establish the operational limits of recalibration under communication drift, we evaluate a mismatched calibration matrix across test latencies $\tau_{\text{test}} \in [0, 60]$ min and calibration horizons $\tau_{\text{cal}} \in [0, 60]$ min (Supplementary Table~A11). Recalibrating at $\tau_{\text{cal}} = 10$ min withstands latency drift up to $\tau_{\text{test}} = 20$ min while maintaining target-satisfying tail coverage (8.2%–9.1% violation). However, recalibration breaks down when staleness surges to $\tau_{\text{test}} = 60$ min (13.8% violation). This demonstrates that under full channel observability, distribution drift does not require neural representation learning: state-conditional recalibration recovers 55.3% of the observed shortage loss (127.6 $\rightarrow$ 57.0 MWh) without updating network weights.

C. Regime 3: Critical State Unobservable $\rightarrow$ Consequence Representation and Decision Boundaries (RQ2)

The targeted utility of machine learning emerges when blade-pitch telemetry is withheld across aggregator boundaries (no_pitch, Table~III). Under pitch unobservability, physical rules and recalibration schemes lose direct state awareness, inflating recalibrated screening penalties from 824,393 kW · h at $\tau = 0$ to 1,378,900 kW · h at $\tau = 60$ min.

Under this regime, the control register is unobservable at deployment, while its electromechanical consequences ($Z_t \rightarrow P_t, Q_t$, trajectory) remain partially recorded in arrived SCA-DA feeds. Learned representations (STGQ-Modular) reconstruct operating states from secondary consequences under **offline privileged supervision**. Specifically, we provide an

TABLE III

MULTIDIMENSIONAL OPERATIONAL PROGRESSION ACROSS TELEMETRY LATENCY $\tau \in [0, 60]$ MIN AND PITCH OBSERVABILITY ON BOUNDARY-ACTIVE REGIMES ($|v_{\text{phys}} - 10.5| \leq 1.0$ m/s, $N = 13,883$ CELLS, $h = 1$, 5-SEED MEAN ON WTB; NOMINAL $\tau = 0$ BASELINE 593,258 kW · h ALIGNS WITHIN 0.6% OF THE PRIMARY FROZEN 589,535 kW · h BOUNDARY BENCHMARK). PENALIZED RESERVE ENERGY IN kW · h ($\rho = 10$).

τ (min)	Pitch	Clean Phys. Metric (Viol.)	Recal. Phys. Metric (Viol.)	STGQ-Mod. Metric (Viol.)	STGQ-Rout. Metric (Viol.)	Operational Regime
0	All	593,258 (7.3%)	593,258 (7.3%)	663,347 (6.1%)	663,915 (7.3%)	Regime I (Lowest Cost)
0	None	961,024 (3.2%)	824,393 (5.4%)	755,794 (5.8%)	756,257 (6.4%)	Regime II (Repr. Adv.)
10	All	753,606 (11.2% [†])	742,211 (7.4%)	786,366 (6.7%)	789,094 (7.6%)	Regime I (Recal. Suff.)
10	None	1,008,743 (4.8%)	928,028 (6.2%)	866,753 (6.1%)	874,156 (6.6%)	Regime II (Repr. Adv.)
20	All	896,198 (14.0% [†])	841,734 (7.1%)	875,838 (7.4%)	879,750 (8.0%)	Regime I (Recal. Suff.)
20	None	1,062,646 (6.1%)	1,010,530 (6.2%)	947,869 (6.7%)	959,934 (7.1%)	Regime II (Repr. Adv.)
30	All	1,069,266 (16.6% [†])	939,190 (7.5%)	987,760 (7.1%)	990,091 (7.6%)	Regime I (Recal. Suff.)
30	None	1,138,774 (7.5%)	1,101,292 (6.7%)	1,054,592 (6.4%)	1,064,238 (6.7%)	Regime II (Repr. Adv.)
60	All	1,656,285 (24.0% [†])	1,228,609 (9.5%)	1,283,855 (9.6%)	1,274,047 (9.6%)	Regime III (Tail Risk)
60	None	1,437,280 (11.8% [†])	1,378,900 (9.0%)	1,330,318 (8.7%)	1,331,840 (8.8%)	Regime II (Repr. Adv.)

[†]Exceeds nominal 10% violation target ($q^* = 0.90$). Bold indicates lowest target-satisfying surrogate screening metric. Under $\tau = 60$ min, target satisfaction is 3/5 seeds (60%) across models.

explicit disclosure that pitch is available in training inputs but withheld in validation/test deployment: blade-pitch registers (β) were present in historical training buffers (x_{hist} channels 7–8) and anchor physics, while pitch was strictly zero-masked and withheld during validation and testing. Trained using historically available pitch information that is withheld at deployment, offline privileged supervision ($\mathcal{L}_{\text{align}}$ cross-entropy on aerodynamic regime labels Z) forces secondary consequence channels to reconstruct the latent boundary space. A clean matched-backbone ablation ($\lambda_{\text{align}} = 5000.0$ vs. 0.0, 5 seeds) confirms that regime supervision improves latent-state organization and reduces representation collapse (NMI 0.819 ± 0.088 vs. 0.240 ± 0.150 , paired $\Delta = +0.579 \pm 0.167$, $p = 0.0015$, 95% CI $[+0.371, +0.787]$; ARI 0.889 ± 0.057 vs. 0.198 ± 0.139 , paired $\Delta = +0.692 \pm 0.153$, $p = 0.0005$, 95% CI $[+0.502, +0.882]$). For posterior calibration, supervision improves Brier score (0.020 ± 0.002 vs. 0.241 ± 0.419 , paired $\Delta = -0.221 \pm 0.418$), though sensitive to representation collapse in Seed 204 (Brier 0.981; excluding Seed 204 yields $\Delta = -0.036 \pm 0.073$, 5-seed 95% CI $[-0.740, +0.299]$, $p = 0.303$). Downstream reserve decision value is evaluated separately, as privileged supervision alone yields an unestablished cost delta ($-311,957 \pm 1,085,873$ kW · h plant-wide, $p = 0.556$; $+15,261 \pm 27,504$ kW · h on the boundary slice, $p = 0.283$; Table~A6c).

Empirical identifiability audits (Table~A7b and Registry) disentangle the recovery mechanism: (1) *Active power signature*: instantaneous power $P_{\text{atv},t}$ provides primary state resolution (Condition A: 59.7% precision, 63.1% recall, F1 0.613, FPR 0.8%, AUROC 0.988), while historical trajectory $P_{\text{atv},t-1}$ retains 40.3% recall and lag-1 $P_{\text{atv},t-1}$ achieves 58.8%, confirming genuine dynamical tracking beyond instantaneous readout; (2) *Consequence suite without current active power*: withholding contemporaneous active power (Condition D) shifts the detector to an asymmetric, high-recall operating point (82.2% recall, 33.4% precision, F1 0.475, FPR 3.2%, AUROC 0.975), confirming that the remaining non-active-power consequence suite retains some collectively recoverable state information; and (3) *Thermal invariance*: conditional ablation of thermal channels ($E_{\text{tmp}}, I_{\text{tmp}}$) yields negligible discrimination change

($\Delta\text{AUROC} = +0.006$, $\Delta\text{F1} = -0.010$, $\Delta\text{Brier} = -0.0008$), confirming that slow thermal dynamics fail to contribute incremental high-frequency boundary information for 10-minute dispatch.

As detailed in Table~III, STGQ-Modular achieves 755,794 kW · h at $\tau = 0$ (an 8.3% saving of 68,599 kW · h over recalibrated physics) and maintains consistent reserve screening reductions of 46.7k–68.6k kW · h across all latencies, doubling transition-window recall (0.416 vs. 0.196).

Preserving the Negative Result (Recoverability $\nRightarrow$ Decision superiority): Crucially, state recoverability does not imply universal decision superiority. Operational validation separates three distinct claims: (i) *Operating-state discrimination* (AUROC > 0.94 , NMI 0.461, ARI 0.647); (ii) *Probabilistic calibration of $P(Z \mid \mathcal{T}_t^{(d)})$* (bulk ECE $\leq 3.46\%$, satisfying the 5% screening gate; tail MCE ≈ 0.75 –0.85); and (iii) *Reserve-decision value*. Against a strong observable wind-speed-conditioned quantile baseline (10 bins, 13.78M kW · h plant-wide), posterior conditioning does not reduce plant-wide reserve screening costs (14.31M kW · h, defined canonically as $\Delta L = L_{\text{posterior}} - L_{\text{wspd}} = +523,044$ kW · h penalty, $p = 0.85$). In steady operation (63.7% of test time), wind-speed binning serves as the operational minimum sufficient model, outperforming the posterior by +410,340 kW · h ($p = 0.002$). Thus, strong wind-speed-conditioned quantiles remain lower-cost plant-wide.

The relative risk-cost trade-off becomes materially more favorable near operating transitions (± 3 steps, 36.3% of records): posterior conditioning selectively reduces violation (7.24% vs. 8.36%) and shortage exposure (86.6k vs. 90.2k kW · h), but this hedge remains more expensive under $\rho = 10$ than wind-speed-binned quantile calibration ($\Delta = +112,704$ kW · h, driven by +149,092 kW · h greater reserve procurement). The significant paired day-level cluster bootstrap interaction ($\Delta_{\text{trans}} - \Delta_{\text{steady}} = -296,123$ kW · h, $p < 0.005$) establishes heterogeneity of the relative cost gap rather than a positive posterior cost advantage. Transition windows are the regions where recovered latent-state information operates as a localized risk-hedging mechanism that contracts tail exposure at added

reserve expense. While transition windows capture 48.4% of savings relative to an unconditioned global baseline, this unconditioned fraction reflects high baseline loss concentration (33.0%) alongside model sensitivity; relative to wind-speed bins, the posterior acts as a localized risk hedge. A deployable validation-frozen hybrid policy matches wind-speed binning plant-wide (13.79M vs. 13.78M kW · h, $p > 0.40$), confirming that machine learning provides conditional regime risk hedging rather than universal cost dominance.

D. Supporting Evidence and Falsification of Alternative Explanations (RQ3)

To ensure that the three-regime principle does not reflect architectural coincidence, we evaluate four sets of diagnostic probes:

1) Architectural Routing Falsification (MoE): Across all evaluated conditions, dynamic Mixture-of-Experts (MoE) routing confers no statistical advantage over unrouted dense baselines. At $h = 1$, STGQ-Dense and STGQ-Routed achieve statistical parity (690,614 vs. 660,990 kW · h, ratio 1.045; seed-paired difference +29,624 kW · h, 95% CI $[-53,774, +113,023]$, $p = 0.380$). At $h = 6$, the paired difference +32,154 kW · h likewise crosses zero ($[-1,148, +65,456]$, Table~II). In multi-horizon dispatch screening across seeds 201–205 (dense_pricing_by_seed.csv), the 5-seed mean surrogate cost difference between soft-gate and soft-dense binning is $< 0.07\%$ (16.065M vs. 16.076M kW · h). Meanwhile, decoupled modular architectures (STGQ-Modular) match or exceed tail reliability under Delay-6 at $h = 1$ (1,284,098 $\pm$ 114,629 kW · h, 9.7% $\pm$ 1.1% violation vs. 1,412,321 kW · h, 10.6% $\pm$ 1.1%). In point forecasting, the boundary-aware model intentionally trades whole-sample RMSE (236.13 $\pm$ 8.41 vs. 224.34 $\pm$ 2.23 for iTransformer and 225.74 $\pm$ 2.60 for Graph WaveNet) to restrict unhedged tail shortfall exposure (53.9 vs. 78.3–82.5 MWh, a 31%–35% reduction; Supplementary Section S3). Operational gains derive from consequence representation and calibrated quantile sizing, not expert routing.

2) Decoupling of Classification Accuracy and Reserve Value: Benchmarking an independent sequence classifier (Graph WaveNet coupled to a live-anchor classifier) against joint representation models reveals a sharp decoupling between classification accuracy and reserve screening cost. Under nominal telemetry, the independent classifier achieves perfect transition recall (1.000) and high precision (0.879), surpassing the joint router (0.960 recall, 0.758 precision). However, in pre-dispatch reserve screening, the independent classifier incurs 86,536,321 kW · h (9.54% violation), suffering a +2.22M kW · h cost penalty over physical bins (84,314,353 kW · h, paired $p = 0.0013$, exact Wilcoxon $p = 0.0625$), and +1.96M kW · h over the joint router (84,577,217 kW · h, CI $[-8.18\text{M}, +12.71\text{M}]$; Supplementary Table~A10c). Discrete classification errors concentrate during steep wind ramps where residuals are largest, compounding Newsvendor penalties.

3) Failure of Heuristic Selective Abstention: Benchmarking selective decision abstention against random rejection and uniform reserve margin expansion (Supplementary Table~A10f) reveals that heuristic uncertainty scores ($U = \Delta r + 50H(\tau)$)

fail under multi-step latency: selective refusal at confidence threshold $c = 0.90$ yields a fleet violation rate of 9.44%, statistically indistinguishable from random rejection (9.50%, $p > 0.40$). Uniform reserve margin expansion achieves an 8.29% violation rate while saving 37,823 kW · h, strictly Pareto-dominating heuristic refusal.

4) Cross-Farm Boundary Probes and Information Heterogeneity: External evaluations across four commercial wind farms reveal heterogeneous information regimes rather than uniform transfer superiority (Supplementary Table~A8). Under local retraining, WTB (134 units, $-38\text{k kW} \cdot \text{h}$) and Penmanshiel (14 operational units, numbered WT01–WT15 excluding WT03, $-2.14\text{M kW} \cdot \text{h}$, both $p < 0.05$) show significant surrogate reductions relative to an unconditioned global quantile baseline, whereas Kelmarsh (6 units) is neutral ($-40\text{k kW} \cdot \text{h}$, 95% CI crossing zero) and LHB (4 units) exhibits a negative transfer/overfitting boundary ($+43\text{k kW} \cdot \text{h}$ penalty, $[-28.7\text{k}, +141.2\text{k}]$). The observed pattern is consistent with exploitable spatial redundancy as a possible moderator (not turbine count alone), while turbine count alone does not explain the variation. Because external comparisons evaluate unconditioned global quantiles without matched external no-graph ablations, these differences are not causally attributed to wake modeling alone. Furthermore, pronounced directional transfer asymmetry (Kelmarsh $\rightarrow$ Penmanshiel NMI 0.77 vs. Penmanshiel $\rightarrow$ Kelmarsh NMI 0.34) supports site-specific retraining or recalibration rather than assuming reliable zero-shot transfer.

5) Point of Common Coupling (PCC) Portfolio Smoothing: Aggregating turbine power flows at the PCC bus ($P_{\text{farm}} = \sum_{i \in \mathcal{T}_t} P_{i,t}$, ~ 121 turbines) unlocks spatial diversification: turbulence cancels across the array, leaving boundary transitions as primary reserve risk. Diversification concentrates in transitional regimes (10%–90% pitching), saving $-1.33\text{M kW} \cdot \text{h}$ (CI $[-1.68\text{M}, -1.07\text{M}]$) over continuous physical pitch; longitudinal records (Kelmarsh, 9 years; Penmanshiel, 8.6 years) confirm operational recalibration under drift (Supplementary Section S4).

VI. DISCUSSION: MINIMUM SUFFICIENT MODEL COMPLEXITY

The empirical findings establish an operational principle of **Minimum Sufficient Model Complexity**: machine learning is needed when observability fails, not whenever prediction is difficult. Dispatch operations should deploy the least complex model sufficient for the arriving information regime:

1) Fresh and Observable $\rightarrow$ Physics: Deterministic aerodynamic curves evaluate pristine telemetry ($\tau = 0$) with asymptotic power conservation ($\partial P / \partial v \propto v^2$ below rated speed, $P \equiv P_{\text{rated}}$ above rated speed). Avoiding out-of-distribution neural variance, they achieve the lowest evaluated surrogate reserve screening cost within the boundary band ($|v_{\text{phys}} - 10.5| \leq 1.0$ m/s, $N = 13,883$ events: 589.5k–881.4k kW · h, target-satisfying $\sim 6.8\%$ violation; Tables~I and~III), rendering deep representations superfluous.

2) Stale but Observable $\rightarrow$ Recalibration: Under observable operational latency (10–30 min), state-conditional recalibra-

tion absorbs 55.3% of shortage loss (127.6 $\rightarrow$ 57.0 MWh) and maintains violation rates below the nominal 10% target (7.1%–7.5%) without updating neural weights. Static recalibration breaks down only under contingency horizons ($\tau = 60$ min, Table~A111, 13.8% violation), which acts as a synthetic reliability breakdown stress endpoint rather than an industrial prevalence claim.

3) *Critical State Unobservable $\rightarrow$ Consequence Representation*: Decisive utility emerges when blade-pitch telemetry is withheld (no_pitch), where physical rules inflate penalties to 1,378,900 kW $\cdot$ h. The spatio-temporal encoder reconstructs latent regimes from secondary electromechanical consequence channels ($Z_t \rightarrow P_t, Q_t$, trajectory, 0.302–0.674 NMI), saving 46.7k–68.6k kW $\cdot$ h.

4) *Boundary Condition (Recoverability $\nRightarrow$ Decision superiority)*: Posterior conditioning does not reduce plant-wide reserve procurement costs over strong wind-speed-binned quantiles ($\Delta L = L_{\text{posterior}} - L_{\text{wspd}} = +523,044$ kW $\cdot$ h penalty, $p = 0.85$). While direct wind-speed binning serves as the plant-wide and steady-state minimum sufficient policy (+410,340 kW $\cdot$ h steady advantage, $p = 0.002$), the relative risk-cost trade-off becomes materially more favorable near operating transitions (± 3 steps): posterior conditioning selectively reduces violation (7.24% vs. 8.36%) and shortage exposure (86.6k vs. 90.2k kW $\cdot$ h), but this hedge remains more expensive under $\rho = 10$ than wind-speed binning ($\Delta = +112,704$ kW $\cdot$ h). The significant interaction contrast ($\Delta_{\text{trans}} - \Delta_{\text{steady}} = -296,123$ kW $\cdot$ h, $p < 0.005$) establishes heterogeneity of the relative cost gap rather than a positive posterior cost advantage. A deployable hybrid policy remains statistically equivalent to wind-speed binning (13.79M vs. 13.78M kW $\cdot$ h, $p > 0.40$) and must not be presented as an improvement. Dynamic MoE routing provides no measurable advantage over unrouted dense baselines (ratio 1.045, $p = 0.380$; seed delta $< 0.07\%$), while modular regression (STGQ-Modular) matches tail reliability (9.7% violation).

5) *Industrial Deployment Guardrails*: (i) *External heterogeneity*: external evaluations across four commercial wind farms reveal heterogeneous information regimes rather than uniform transfer superiority (Supplementary Table~A8). STGQ reduces surrogate cost relative to global quantiles on WTB and Penmanshiel ($p < 0.05$), while Kelmarsh is neutral and LHB exhibits a negative transfer/overfitting boundary (+43k kW $\cdot$ h penalty), consistent with limited exploitable spatial redundancy as a possible moderator (not turbine count alone), while turbine count alone does not explain the variation; (ii) *Local retraining*: directional transfer asymmetry (Kelmarsh $\rightarrow$ Penmanshiel NMI 0.77 vs. 0.34) supports site-specific retraining or recalibration rather than assuming reliable zero-shot transfer; (iii) *Abstention limits*: selective refusal fails under latency ($p > 0.40$), favoring uniform margin expansion.

VII. LIMITATIONS

The empirical findings are bounded by four operational and experimental scope conditions:

1) *Synthetic Stress*: Impairments are evaluated under synthetic stress (10–30 min backlogs, 60-min latency, Gilbert-Elliott dropouts); field units may exhibit sensor icing or individual pitch actions outside supervisory records.

2) *Risk Surrogacy*: Evaluation operates at Level-1 pre-dispatch operating reserve screening via the PSREI newsvendor surrogate ($\rho = 10$), abstracting Level-2 bulk transmission AC-OPF and balancing settlements.

3) *Horizon Scope*: Benchmarks use a 35-day continuous test split across five random seeds (201–205), omitting multi-annual climatological cycles. Rated wind speeds vary by model (12.5 m/s on MM92 vs. 14.5 m/s on MM82), requiring OEM calibration.

4) *Deployment Boundaries*: Cross-site STGQ effects relative to the global quantile baseline remain heterogeneous, consistent with exploitable spatial redundancy as a possible moderator rather than turbine count alone; LHB exhibits a negative transfer/overfitting boundary, consistent with limited exploitable spatial redundancy and site-specific heterogeneity (+43k kW $\cdot$ h penalty). Because external benchmarks (Supplementary Table~A8) evaluate unconditioned global quantiles without matched external no-graph controls, wake modeling is not causally isolated as the sole mechanism for external gains. Directional transfer asymmetry (Kelmarsh $\rightarrow$ Penmanshiel NMI 0.77 vs. Penmanshiel $\rightarrow$ Kelmarsh NMI 0.34) supports site-specific retraining or recalibration rather than assuming reliable zero-shot transfer. The study assumes historical blade-pitch availability during offline training, followed by pitch withholding at deployment. Performance in settings where pitch was never recorded remains untested.

VIII. CONCLUSION

This study repositions machine learning in reserve screening from an indiscriminate forecasting tool to a targeted conditional inference mechanism governed by SCADA observability regimes. The governing principle is: *use the least complex model sufficient for the current information regime; machine learning is needed when observability fails, not whenever prediction is difficult*. First, under fresh telemetry, deterministic aerodynamic power curves minimize surrogate reserve penalties within the boundary band ($|v_{\text{phys}} - 10.5| \leq 1.0$ m/s, $N = 13,883$ events: 589,535 kW $\cdot$ h at $h = 1$, 881,367 kW $\cdot$ h at $h = 6$, $\sim 6.8\%$ violation), collapsing only under unadapted latency (24.0% violation). Second, under observable delays, state-conditional recalibration absorbs 55.3% of shortage loss (127.6 $\rightarrow$ 57.0 MWh) without updating neural weights. Third, when blade pitch is withheld, learned representations infer latent regimes from secondary electromechanical consequences ($Z_t \rightarrow P_t, Q_t$, trajectory), saving 46.7k–68.6k kW $\cdot$ h and doubling transition recall (0.416 vs. 0.196). However, state recoverability does not imply plant-wide decision superiority: strong wind-speed-conditioned quantiles remain lower-cost plant-wide, with posterior representations functioning strictly as transition-window risk hedges. Finally, dynamic MoE routing confers no statistical advantage over unrouted architectures ($p = 0.380$). Dispatch systems should deploy deterministic rules under clean telemetry, recalibration under observable

drift, and learned representations when critical control states are unobservable.

DECLARATIONS

AI Use Statement: The authors used OpenAI ChatGPT and Codex for drafting, formatting, and prose refinement, and for developing audit scripts. All scientific hypotheses, mathematical formulations, experimental designs, data processing, statistical analyses, findings, and final approval remain the sole responsibility of the authors.

Code and Data Availability: Raw SCADA datasets (WTB, ENGIE La Haute Borne, Kelmarsh, Penmanshiel) are public. Code, manifests, derived tables, and checkpoints are provided in the replication repository across declared seeds. The analysis involves no human subjects. # References {unnumbered}

- [1] J. Dowell and P. Pinson, "Very-short-term probabilistic wind power forecasts by sparse vector autoregression," *IEEE Transactions on Smart Grid*, vol. 7, no. 2, pp. 1063–1070, 2016, doi: 10.1109/TS-G.2015.2424078.
- [2] J. Kruse, E. Cramer, B. Sch"after, and D. Witthaut, "Physics-informed machine learning for power grid frequency modeling," *PRX Energy*, vol. 2, no. 4, p. 043003, 2023, doi: 10.1103/PRXEnergy.2.043003.
- [3] P. Pinson, "Wind energy: Forecasting challenges for its operational management," *Statistical Science*, vol. 28, no. 4, pp. 564–585, 2013, doi: 10.1214/13-sts445.
- [4] R. Doherty and M. O'Malley, "A new approach to quantify reserve demand in systems with significant installed wind capacity," *IEEE Transactions on Power Systems*, vol. 20, no. 2, pp. 587–595, 2005, doi: 10.1109/TPWRS.2005.846206.
- [5] Y. Wang, F. Zhang, H. Kou, R. Zou, Q. Hu, and J. Wang, "A review of predictive uncertainty modeling techniques and evaluation metrics in probabilistic wind speed and wind power forecasting," *Applied Energy*, vol. 396, p. 126234, 2025, doi: 10.1016/j.apenergy.2025.126234.
- [6] J. G. Slootweg, S. W. H. de Haan, H. Polinder, and W. L. Kling, "General model for representing variable speed wind turbines in power system dynamics simulations," *IEEE Transactions on Power Systems*, vol. 18, no. 1, pp. 144–151, 2003, doi: 10.1109/TPWRS.2002.807113.
- [7] E. Gaertner *et al.*, "Definition of the IEA wind 15-megawatt offshore reference wind turbine," National Renewable Energy Laboratory (NREL), Golden, CO, NREL/TP-5000-75698, 2020. doi: 10.2172/1603478.
- [8] J. Tautz-Weinert and S. J. Watson, "Using SCADA data for wind turbine condition monitoring: A review," *IET Renewable Power Generation*, vol. 11, no. 4, pp. 382–394, 2017, doi: 10.1049/iet-rpg.2016.0248.
- [9] M. A. Ullah, K. Mikhaylov, and H. Alves, "Enabling mMTC in remote areas: LoRaWAN and LEO satellite integration for offshore wind farm monitoring," *IEEE Transactions on Industrial Informatics*, vol. 18, no. 6, pp. 3744–3753, 2022, doi: 10.1109/TII.2021.3112386.
- [10] S. Daenens, T. Verstraeten, P.-J. Daems, A. Nowé, and J. Helsens, "Spatio-temporal graph neural networks for power prediction in offshore wind farms using SCADA data," *Wind Energy Science*, vol. 10, no. 6, pp. 1137–1152, 2025, doi: 10.5194/wes-10-1137-2025.
- [11] Z. Wu, S. Pan, G. Long, J. Jiang, and C. Zhang, "Graph WaveNet for deep spatial-temporal graph modeling," in *Proceedings of the 28th international joint conference on artificial intelligence (IJCAI)*, 2019, pp. 1907–1913. doi: 10.24963/ijcai.2019/264.
- [12] S. Guo, Y. Lin, N. Feng, C. Song, and H. Wan, "Attention based spatial-temporal graph convolutional networks for traffic flow forecasting," in *Proceedings of the AAAI conference on artificial intelligence*, 2019, pp. 922–929. doi: 10.1609/aaai.v33i01.3301898.
- [13] L. Bai, L. Yao, C. Li, X. Wang, and C. Wang, "Adaptive graph convolutional recurrent network for traffic forecasting," in *Advances in neural information processing systems*, 2020, pp. 17804–17815.
- [14] J. Park and J. Park, "Physics-induced graph neural network: An application to wind-farm power estimation," *Energy*, vol. 187, p. 115883, 2019, doi: 10.1016/j.energy.2019.115883.
- [15] D. Kim, G. Ryu, and C. Moon, "Accuracy of a short-term wind power forecasting model based on deep learning using LiDAR-SCADA integration: A case study of the 400-MW Anholt offshore wind farm," *Applied Energy*, vol. 373, p. 123882, 2024, doi: 10.1016/j.apenergy.2024.123882.
- [16] N. Zehtabian-Rezaie, A. Iosifidis, and M. Abkar, "Physics-guided machine learning for wind-farm power prediction: Toward interpretability and generalizability," *PRX Energy*, vol. 2, no. 1, p. 013009, 2023, doi: 10.1103/PRXEnergy.2.013009.
- [17] E. A. Bossanyi, "The design of closed loop controllers for wind turbines," *Wind Energy*, vol. 3, no. 3, pp. 149–163, 2000, doi: 10.1002/we.34.
- [18] F. D. Bianchi, H. De Battista, and R. J. Mantz, *Wind turbine control systems: Principles, modelling and gain scheduling design*. London: Springer, 2006. doi: 10.1007/1-84628-493-7.
- [19] L. Y. Pao and K. E. Johnson, "Control of wind turbines: Approaches, challenges, and recent developments," *IEEE Control Systems Magazine*, vol. 31, no. 2, pp. 44–62, 2011, doi: 10.1109/MCS.2010.939962.
- [20] C. Zhao *et al.*, "A physics-aware dynamic graph and mixture-of-experts framework for wind power forecasting," *International Journal of Electrical Power & Energy Systems*, vol. 176, p. 111769, 2026, doi: 10.1016/j.ijepes.2026.111769.
- [21] J. B. Bremnes, "Probabilistic wind power forecasts using local quantile regression," *Wind Energy*, vol. 7, no. 1, pp. 47–54, 2004, doi: 10.1002/we.107.
- [22] H. A. Nielsen, H. Madsen, and T. S. Nielsen, "Using quantile regression to extend an existing wind power forecasting system with probabilistic forecasts," *Wind Energy*, vol. 9, no. 1–2, pp. 95–108, 2006, doi: 10.1002/we.180.
- [23] N. Shazeer *et al.*, "Outrageously large neural networks: The sparsely-gated mixture-of-experts layer," in *International conference on learning representations*, 2017.
- [24] W. Fedus, B. Zoph, and N. Shazeer, "Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity," *Journal of Machine Learning Research*, vol. 23, no. 120, pp. 1–39, 2022.
- [25] J. Zhou *et al.*, "SDWPF: A dataset for spatial dynamic wind power forecasting over a large turbine array," *Scientific Data*, vol. 11, p. 649, 2024, doi: 10.1038/s41597-024-03427-5.
- [26] H. Hersbach *et al.*, "The ERA5 global reanalysis," *Quarterly Journal of the Royal Meteorological Society*, vol. 146, no. 730, pp. 1999–2049, 2020, doi: 10.1002/qj.3803.
- [27] B. J. Pierre *et al.*, "Design of the Pacific DC Intertie wide area damping controller," *IEEE Transactions on Power Systems*, vol. 34, no. 5, pp. 3594–3604, 2019, doi: 10.1109/TPWRS.2019.2903782.
- [28] G. Ravikumar and M. Govindarasu, "Anomaly detection and mitigation for wide-area damping control using machine learning," *IEEE Transactions on Smart Grid*, vol. 15, no. 6, pp. 5939–5951, 2024, doi: 10.1109/TSG.2020.2995313.